

customer-churn-project

April 24, 2025

1 PROJECT OVERVIEW

Customer churn, or customer attrition, refers to when a customer ceases their relationship with a company or service provider. In today's highly competitive business environment, retaining customers is a critical factor for long-term success. Predicting customer churn can help organizations take proactive steps to retain customers, thus minimizing revenue loss. This project aims to build a machine learning model that can predict whether a customer will churn based on their demographic, account, and service-related data.

2 PROBLEM STATEMENT

The goal of this project is to develop a classification model that predicts whether a customer will churn. Using demographic data (such as gender, senior citizen status, and tenure), along with information about the services they use (such as internet service, phone service, and online security), we will attempt to build a model that helps the company identify customers who are at a high risk of churning. By predicting customer churn, the company can proactively design retention strategies to keep these customers, thereby improving customer satisfaction and reducing financial loss.

```
[22]: # Importing useful libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import \
    accuracy_score, precision_score, recall_score, f1_score, classification_report, confusion_matrix

# 1. Load Dataset
df = pd.read_csv("C:/Users/Himanshu/Downloads/customer_data.csv")
df.head()
print("\nDataset Info:")
df.info()
print("\nMissing Values:")
(df.isnull().sum())
```

Dataset Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 7043 entries, 0 to 7042

Data columns (total 21 columns):

#	Column	Non-Null Count	Dtype
0	customerID	7043 non-null	object
1	gender	7043 non-null	object
2	SeniorCitizen	7043 non-null	int64
3	Partner	7043 non-null	object
4	Dependents	7043 non-null	object
5	tenure	7043 non-null	int64
6	PhoneService	7043 non-null	object
7	MultipleLines	7043 non-null	object
8	InternetService	7043 non-null	object
9	OnlineSecurity	7043 non-null	object
10	OnlineBackup	7043 non-null	object
11	DeviceProtection	7043 non-null	object
12	TechSupport	7043 non-null	object
13	StreamingTV	7043 non-null	object
14	StreamingMovies	7043 non-null	object
15	Contract	7043 non-null	object
16	PaperlessBilling	7043 non-null	object
17	PaymentMethod	7043 non-null	object
18	MonthlyCharges	7043 non-null	float64
19	TotalCharges	7032 non-null	float64
20	Churn	7043 non-null	object

dtypes: float64(2), int64(2), object(17)

memory usage: 1.1+ MB

Missing Values:

```
[22]: customerID      0
      gender         0
      SeniorCitizen  0
      Partner        0
      Dependents     0
      tenure         0
      PhoneService   0
      MultipleLines  0
      InternetService 0
      OnlineSecurity 0
      OnlineBackup   0
      DeviceProtection 0
      TechSupport    0
      StreamingTV    0
```

```

StreamingMovies      0
Contract              0
PaperlessBilling      0
PaymentMethod         0
MonthlyCharges        0
TotalCharges          11
Churn                 0
dtype: int64

```

```

[54]: df.dropna(inplace=True)
      # Feature Engineering
      # Keeping tenure as a numerical feature instead of binning
      # Encoding Categorical Variables
      label_encoders = {}
      for col in df.select_dtypes(include=['object']).columns:
          le = LabelEncoder()
          df[col] = le.fit_transform(df[col])
          label_encoders[col] = le

      # Feature Scaling
      scaler = StandardScaler()
      feature_columns = df.columns.difference(['customerID', 'Churn'])
      df[feature_columns] = scaler.fit_transform(df[feature_columns])

      # Splitting Data
      X = df.drop(columns=['Churn'])
      y = df['Churn']
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
          ↪2, random_state=42, stratify=y)

      # Handling Class Imbalance
      smote = SMOTE(random_state=42)
      X_train, y_train = smote.fit_resample(X_train, y_train)

```

```

[26]: # Model Training with XGBoost
      xgb = XGBClassifier(objective='binary:logistic',
          ↪eval_metric='logloss', random_state=42)
      param_grid = {
          'n_estimators': [50, 100, 200],
          'max_depth': [3, 6, 9],
          'learning_rate': [0.01, 0.1, 0.2]
      }
      grid_search = GridSearchCV(xgb, param_grid, cv=5, scoring='f1', n_jobs=-1)
      grid_search.fit(X_train, y_train)
      # Best Model
      best_model = grid_search.best_estimator_
      y_pred = best_model.predict(X_test)


```

```
[27]: # Model Evaluation
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred))
print("Recall:", recall_score(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
# Confusion Matrix
plt.figure(figsize=(5,4))
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d', cmap='Blues')
3
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```

Accuracy: 0.767590618336887
Precision: 0.5547785547785548
Recall: 0.6363636363636364
F1 Score: 0.5927770859277709

Classification Report:

	precision	recall	f1-score	support
0	0.86	0.82	0.84	1033
1	0.55	0.64	0.59	374
accuracy			0.77	1407
macro avg	0.71	0.73	0.72	1407
weighted avg	0.78	0.77	0.77	1407



```
[ ]: # Predicting for New Data
def predict_churn(new_data):
    new_df = pd.DataFrame([new_data])
    # Encoding categorical variables
    for col, le in label_encoders.items():
        if col in new_df.columns:
            new_df[col] = new_df[col].apply(lambda x: le.transform([x])[0]
            if x in le.classes_ else -1)
    # Ensuring numerical columns are correctly formatted
    for col in feature_columns:
        if col in new_df.columns:
            new_df[col] = pd.to_numeric(new_df[col], errors='coerce')
    new_df.fillna(0, inplace=True) # Handling NaN values that may arise
    new_df[feature_columns] = scaler.transform(new_df[feature_columns])
    return best_model.predict(new_df)[0]
```

```
[86]: # Example Prediction
new_customer = {'gender': 'Male', 'SeniorCitizen': 0, 'Partner': 'No',
                'Dependents': 'No', 'tenure': 12.0, 'PhoneService': 'Yes',
                'MultipleLines': 'No', 'InternetService': 'Fiber optic', 'OnlineSecurity': 'No',
                'OnlineBackup': 'No',
```

```
'DeviceProtection': 'No', 'TechSupport': 'No', 'StreamingTV':'Yes',  
↪ 'StreamingMovies': 'No',  
'Contract': 'Month-to-month', 'PaperlessBilling': 'Yes', 'PaymentMethod':  
↪ 'Electronic check',  
'MonthlyCharges': 70.35, 'TotalCharges': 850.0}  
print("Predicted Churn (1 = Yes, 0 = No):", predict_churn(new_customer))
```

Predicted Churn (1 = Yes, 0 = No): None

3 4. Insights

High Risk: Month-to-month contracts, high monthly charges, and lack of add-ons.

Retention Strategy: Offer discounts or upgrades for high-risk churn profiles.

Services Impact: OnlineSecurity, TechSupport reduce churn likelihood.

4 Video Explanation:

<https://drive.google.com/file/d/1F1IMVpdOxfPuhkJ5dGNQrkdqd3yZ7Fcj/view?usp=sharing>

[]: